

```
In [1]: import os
```

```
In [2]: import os  
os.chdir("D:/Certificates/Data Science/Edwisor/Project 1")
```

```
In [3]: # check current working directory  
os.getcwd()
```

```
Out[3]: 'D:\\Certificates\\Data Science\\Edwisor\\Project 1'
```

```
In [4]: print(os.listdir(os.getcwd()))  
  
['.RData', '.Rhistory', 'bank-loan.csv', 'DataN0108 (1).pdf', 'In.docx',  
'Project 1.csv', 'project report_.pdf', 'Project.docx', 'Python coding.docx',  
'R Coding.docx', 'R.R']
```

```
In [5]: import numpy as np  
import pandas as pd  
import matplotlib.pyplot as plt  
import seaborn as sns  
sns.set()  
%matplotlib inline  
from sklearn.model_selection import cross_val_score  
from sklearn.model_selection import train_test_split  
from sklearn.ensemble import RandomForestClassifier  
from sklearn.metrics import accuracy_score  
from sklearn.metrics import classification_report  
from sklearn.preprocessing import StandardScaler  
from sklearn.pipeline import make_pipeline  
from sklearn import svm  
from sklearn.preprocessing import scale  
from sklearn.model_selection import GridSearchCV  
from sklearn.linear_model import LogisticRegression  
from sklearn.metrics import precision_recall_curve  
from sklearn.metrics import auc
```

```
from sklearn.metrics import roc_curve
from sklearn.metrics import roc_auc_score
from sklearn.decomposition import PCA
from sklearn.ensemble import GradientBoostingClassifier
```

```
In [6]: bank = pd.read_csv("bank-loan.csv")
```

```
In [7]: bank.columns
```

```
Out[7]: Index(['age', 'ed', 'employ', 'address', 'income', 'debtinc', 'creddeb  
t',  
            'othdebt', 'default'],  
            dtype='object')
```

```
In [8]: bank.head(5)
```

```
Out[8]:
```

	age	ed	employ	address	income	debtinc	creddebt	othdebt	default
0	41	3	17	12	176	9.3	11.359392	5.008608	1.0
1	27	1	10	6	31	17.3	1.362202	4.000798	0.0
2	40	1	15	14	55	5.5	0.856075	2.168925	0.0
3	41	1	15	14	120	2.9	2.658720	0.821280	0.0
4	24	2	2	0	28	17.3	1.787436	3.056564	1.0

```
In [9]: bank.tail(5)
```

```
Out[9]:
```

	age	ed	employ	address	income	debtinc	creddebt	othdebt	default
845	34	1	12	15	32	2.7	0.239328	0.624672	NaN
846	32	2	12	11	116	5.7	4.026708	2.585292	NaN
847	48	1	13	11	38	10.8	0.722304	3.381696	NaN
848	35	2	1	11	24	7.8	0.417456	1.454544	NaN

	age	ed	employ	address	income	debtinc	creddebt	othdebt	default
849	37	1	20	13	41	12.9	0.899130	4.389870	NaN

In [10]: bank.info()

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 850 entries, 0 to 849
Data columns (total 9 columns):
age                850 non-null int64
ed                850 non-null int64
employ            850 non-null int64
address           850 non-null int64
income            850 non-null int64
debtinc           850 non-null float64
creddebt          850 non-null float64
othdebt           850 non-null float64
default           700 non-null float64
dtypes: float64(4), int64(5)
memory usage: 59.9 KB
```

In [11]: bank.shape

Out[11]: (850, 9)

In [12]: bank.describe()


Out[12]:

	age	ed	employ	address	income	debtinc	creddebt	othdebt	default
count	850.000000	850.000000	850.000000	850.000000	850.000000	850.000000	850.000000	850.000000	700.000000
mean	35.029412	1.710588	8.565882	8.371765	46.675294	10.171647	1.576805	3.029412	0.835294
std	8.041432	0.927784	6.777884	6.895016	38.543054	6.719441	2.125840	3.329412	0.371765
min	20.000000	1.000000	0.000000	0.000000	13.000000	0.100000	0.011696	0.000000	0.000000
25%	29.000000	1.000000	3.000000	3.000000	24.000000	5.100000	0.382176	1.000000	0.000000
50%	34.000000	1.000000	7.000000	7.000000	35.000000	8.700000	0.885091	2.000000	0.000000

	age	ed	employ	address	income	debtinc	creddebt	o
75%	41.000000	2.000000	13.000000	12.000000	55.750000	13.800000	1.898440	3.9
max	56.000000	5.000000	33.000000	34.000000	446.000000	41.300000	20.561310	35.1


```
In [13]: ##### univariate analysis and bivariate analysis #####
#####
# analysis for single variable in the dataset and relation between 2 va
riables.
sns.distplot(bank["age"], kde = True , bins = 40)
```

Out[13]: <matplotlib.axes._subplots.AxesSubplot at 0x2797f155a88>




```
In [14]: sns.distplot(bank["income"], kde = True , bins = 40)
```

Out[14]: <matplotlib.axes._subplots.AxesSubplot at 0x2797f4e8e48>



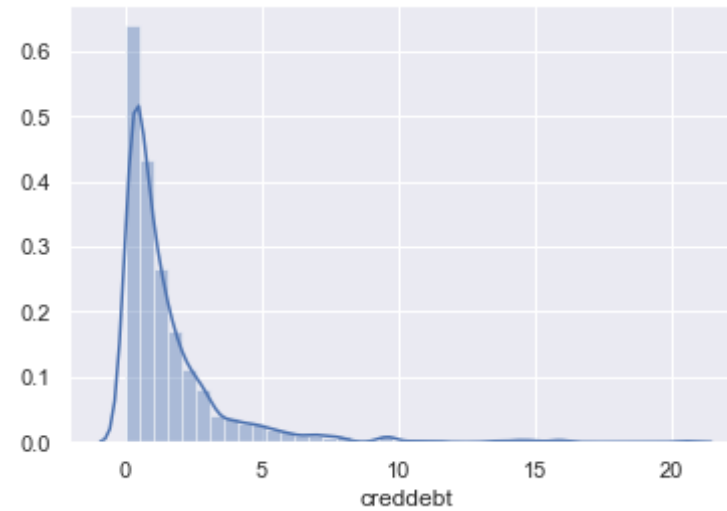
```
In [15]: sns.distplot(bank["othdebt"], kde = True , bins = 40)
```

```
Out[15]: <matplotlib.axes._subplots.AxesSubplot at 0x2797f5c2288>
```




```
In [16]: sns.distplot(bank["creddebt"], kde = True , bins = 40)
```

Out[16]: <matplotlib.axes._subplots.AxesSubplot at 0x2797f5dd548>




```
In [17]: sns.distplot(bank["debtinc"], kde = True , bins = 40)
```

Out[17]: <matplotlib.axes._subplots.AxesSubplot at 0x2797f75e0c8>




```
In [18]: sns.countplot( x = "default" , data = bank)
```

```
Out[18]: <matplotlib.axes._subplots.AxesSubplot at 0x2797f845a48>
```




```
In [19]: sns.countplot( x = "ed" , data = bank, hue = "default")
```

```
Out[19]: <matplotlib.axes._subplots.AxesSubplot at 0x2797f8b9308>
```



```
In [20]: sns.set_style("whitegrid")
plt.figure(figsize = (15,10))
sns.countplot( x = "income" , data = bank, hue = "default")
```

Out[20]: <matplotlib.axes._subplots.AxesSubplot at 0x2797f94ed88>



```
In [21]: ##### missing value analysis #####

#Calculating the null values in the dataframe
missing_value = pd.DataFrame(bank.isnull().sum())
missing_value = (missing_value/len(bank))*100
missing_value.reset_index()

missing_value = missing_value.rename(columns = {'index': 'Variables', 0
```



```

: 'Missing_percentage'})
#Arranging Missing Values in Decreasing Order
missing_value = missing_value.sort_values('Missing_percentage', ascending = False)
#save output results
missing_value.to_csv("Missing_perc.csv", index = False)
missing_value


```

Out[21]:

Missing_percentage	
default	17.647059
age	0.000000
ed	0.000000
employ	0.000000
address	0.000000
income	0.000000
debtinc	0.000000
creddebt	0.000000
othdebt	0.000000

In [22]: `sns.heatmap(bank.isnull(),yticklabels=False,cbar=False,cmap="viridis")`

Out[22]: `<matplotlib.axes._subplots.AxesSubplot at 0x2797fd57508>`




```
In [23]: bank.isnull().sum()
```

```
Out[23]: age          0
         ed           0
         employ       0
         address      0
         income       0
         debtinc      0
         creddebt     0
         othdebt      0
         default     150
         dtype: int64
```

```
In [24]: sns.boxplot( x = "age" , data = bank , orient = "v")
```

```
Out[24]: <matplotlib.axes._subplots.AxesSubplot at 0x2790005b488>
```



```
In [25]: bank.default=bank.default.fillna(2)
```

```
In [26]: bank.default=bank.default.astype(int)
```

```
In [27]: bank.head(2)
```

Out[27]:

	age	ed	employ	address	income	debtinc	creddebt	othdebt	default
0	41	3	17	12	176	9.3	11.359392	5.008608	1
1	27	1	10	6	31	17.3	1.362202	4.000798	0


```
In [28]: bank.tail(2)
```

Out[28]:

	age	ed	employ	address	income	debtinc	creddebt	othdebt	default
848	35	2	1	11	24	7.8	0.417456	1.454544	2
849	37	1	20	13	41	12.9	0.899130	4.389870	2


```
In [29]: df_1 = bank.corr()
sns.heatmap(df_1 , annot = True , cmap = "coolwarm")
```

Out[29]: <matplotlib.axes._subplots.AxesSubplot at 0x279000d3e08>




In [30]: sns.countplot(bank.default)

Out[30]: <matplotlib.axes._subplots.AxesSubplot at 0x2790020d188>



```
In [31]: sns.barplot(x='default',y='income',data=bank)
```

```
Out[31]: <matplotlib.axes._subplots.AxesSubplot at 0x2790026fc08>
```



```
In [32]: ##### selecting all missing values from dataset and we will predict those  
         ## machine learning algorithm.  
         train=bank.loc[bank['default']!=2]  
         print(train.head(5))  
         print(train.tail(5))
```

	age	ed	employ	address	income	debtinc	creddebt	othdebt	default
0	41	3	17	12	176	9.3	11.359392	5.008608	
1	27	1	10	6	31	17.3	1.362202	4.000798	
2	40	1	15	14	55	5.5	0.856075	2.168925	
3	41	1	15	14	120	2.9	2.658720	0.821280	
4	24	2	2	0	28	17.3	1.787436	3.056564	

	age	ed	employ	address	income	debtinc	creddebt	othdebt	def
ault									
695	36	2	6	15	27	4.6	0.262062	0.979938	
1									
696	29	2	6	4	21	11.5	0.369495	2.045505	
0									
697	33	1	15	3	32	7.6	0.491264	1.940736	
0									
698	45	1	19	22	77	8.4	2.302608	4.165392	
0									
699	37	1	12	14	44	14.7	2.994684	3.473316	
0									

```
In [33]: train.default.unique()
```

```
Out[33]: array([1, 0], dtype=int64)
```

```
In [34]: test=bank.loc[bank.default==2]
test=test.iloc[:,0:8]
print(test.head(2))
print(test.tail(2))
```

	age	ed	employ	address	income	debtinc	creddebt	othdebt
700	36	1	16	13	32	10.9	0.544128	2.943872
701	50	1	6	27	21	12.9	1.316574	1.392426
	age	ed	employ	address	income	debtinc	creddebt	othdebt
848	35	2	1	11	24	7.8	0.417456	1.454544
849	37	1	20	13	41	12.9	0.899130	4.389870

```
In [35]: X =train[['age', 'ed', 'employ', 'address', 'income', 'debtinc', 'creddebt',
'othdebt']]
y = train['default']
```

```
In [36]: Xtrain, Xtest, ytrain, ytest = train_test_split(X, y, random_state=42,
test_size=0.25)
```

```
In [37]: from sklearn.tree import DecisionTreeClassifier
```

```
In [38]: clf_gini = DecisionTreeClassifier(criterion = "gini", random_state = 10
,
max_depth=5, min_samples_leaf=7)
clf_gini.fit(Xtrain, ytrain)
```

```
Out[38]: DecisionTreeClassifier(class_weight=None, criterion='gini', max_depth=
5,
max_features=None, max_leaf_nodes=None,
min_impurity_decrease=0.0, min_impurity_split=None,
min_samples_leaf=7, min_samples_split=2,
min_weight_fraction_leaf=0.0, presort=False,
random_state=10, splitter='best')
```

```
In [39]: dt_pred = clf_gini.predict(Xtest)
```

```
In [40]: dt_pred1 = clf_gini.predict_proba(Xtest)[:,:1]
```

```
In [41]: accuracy_score(ytest, dt_pred)
```

```
Out[41]: 0.7828571428571428
```

```
In [42]: print(classification_report(ytest, dt_pred))
```

	precision	recall	f1-score	support
0	0.85	0.87	0.86	132
1	0.56	0.51	0.54	43
accuracy			0.78	175
macro avg	0.70	0.69	0.70	175
weighted avg	0.78	0.78	0.78	175

```
In [43]: precision_dc, recall_dc, thresholds_dc = precision_recall_curve(ytest,
```

```
dt_pred1)
```

```
In [44]: fpr_dc, tpr_dc, thresholds_dc = roc_curve(ytest, dt_pred1)
```

```
In [45]: ##### Random Forest#####  
clf_rf = RandomForestClassifier(random_state=42)
```

```
In [46]: clf_rf.fit(Xtrain, ytrain)
```

```
C:\Users\chinnababu\Anaconda3\lib\site-packages\sklearn\ensemble\forests.py:245: FutureWarning: The default value of n_estimators will change from 10 in version 0.20 to 100 in 0.22.  
"10 in version 0.20 to 100 in 0.22.", FutureWarning)
```

```
Out[46]: RandomForestClassifier(bootstrap=True, class_weight=None, criterion='gini',  
                                max_depth=None, max_features='auto', max_leaf_nodes=None,  
                                min_impurity_decrease=0.0, min_impurity_split=None,  
                                min_samples_leaf=1, min_samples_split=2,  
                                min_weight_fraction_leaf=0.0, n_estimators=10,  
                                n_jobs=None, oob_score=False, random_state=42, verbose=0,  
                                warm_start=False)
```

```
In [47]: y_pred_rf = clf_rf.predict(Xtest)
```

```
In [48]: cv_scores = cross_val_score(clf_rf, Xtrain, ytrain, cv = 5)  
print("Average 5-Fold CV Score: {}".format(np.mean(cv_scores)))
```

```
Average 5-Fold CV Score: 0.7714285714285715
```

```
In [49]: cv_scores = cross_val_score(clf_rf, Xtrain, ytrain, cv = 5, scoring =  
    'roc_auc')  
print("Average 5-Fold CV Score using ROC scoring: {}".format(np.mean(cv_scores)))
```


Average 5-Fold CV Score using ROC scoring: 0.7528293135435994

```
In [50]: accuracy_score(ytest, y_pred_rf)
```

```
Out[50]: 0.7542857142857143
```

```
In [51]: n_space = np.array([5, 6, 10, 12, 15, 50, 100, 200, 500])
criterion_vals = ['gini', 'entropy']
max_features_vals = ['auto', 'sqrt', 'log2']
min_samples_leaf_sp = [1,5,10,25,50]
bootstrap_sp = [True, False]

param_grid = {'n_estimators': n_space, 'criterion' : criterion_vals,
              'max_features':max_features_vals, 'min_samples_leaf': min
              _samples_leaf_sp,
              'bootstrap': bootstrap_sp}
```

```
In [52]: rf_clf_tuning = GridSearchCV(clf_rf, param_grid, cv=5)
```

```
In [53]: rf_clf_tuning.fit(Xtrain, ytrain)
```

```
Out[53]: GridSearchCV(cv=5, error_score='raise-deprecating',
                      estimator=RandomForestClassifier(bootstrap=True, class_weight=None,
                                                         criterion='gini', max_depth=None,
                                                         max_features='auto',
                                                         max_leaf_nodes=None,
                                                         min_impurity_decrease=0.0,
                                                         min_impurity_split=None,
                                                         min_samples_leaf=1,
                                                         min_samples_split=2,
                                                         min_weight_fraction_leaf=0.0,
                                                         n_estimators=10, n_jobs=1,
                                                         random_state=None,
                                                         verbose=0, warm_start=False),
                      param_grid={'bootstrap': [True, False],
                                   'criterion': ['gini', 'entropy'],
                                   'max_features': ['auto', 'sqrt', 'log2'],
                                   'min_samples_leaf': [1, 5, 10, 25, 50],
                                   'n_estimators': [5, 6, 10, 12, 15, 50, 100, 200, 500]},
                      scoring='roc_auc',
                      verbose=0)
```

```

tate=42,
se),
        iid='warn', n_jobs=None,
        param_grid={'bootstrap': [True, False],
                    'criterion': ['gini', 'entropy'],
                    'max_features': ['auto', 'sqrt', 'log2'],
                    'min_samples_leaf': [1, 5, 10, 25, 50],
                    'n_estimators': array([ 5,  6, 10, 12, 1
5,  50, 100, 200, 500])},
        pre_dispatch='2*n_jobs', refit=True, return_train_score=False,
        scoring=None, verbose=0)

```

```

In [54]: print("Tuned RF Parameters: {}".format(rf_clf_tuning.best_params_))
print("Best score is {}".format(rf_clf_tuning.best_score_))

```

```

Tuned RF Parameters: {'bootstrap': True, 'criterion': 'gini', 'max_features': 'log2', 'min_samples_leaf': 5, 'n_estimators': 200}
Best score is 0.8019047619047619

```

```

In [64]: best_rf_clf = RandomForestClassifier(criterion = 'gini', bootstrap = True,
                                             max_features = 'log2', min_samples_leaf = 5, n_estimators = 200)

```

```

In [65]: best_rf_clf.fit(Xtrain, ytrain)

```

```

Out[65]: RandomForestClassifier(bootstrap=True, class_weight=None, criterion='gini',
                                max_depth=None, max_features='log2', max_leaf_nodes=None,
                                min_impurity_decrease=0.0, min_impurity_split=None,
                                min_samples_leaf=5, min_samples_split=2,
                                min_weight_fraction_leaf=0.0, n_estimators=200,

```

```
n_jobs=None, oob_score=False, random_state=None,  
verbose=0, warm_start=False)
```

```
In [66]: y_best_rf_preds = best_rf_clf.predict(Xtest)
```

```
In [67]: cv_scores = cross_val_score(best_rf_clf, Xtrain, ytrain, cv = 5)  
print("Average 5-Fold CV Score: {}".format(np.mean(cv_scores)))
```

Average 5-Fold CV Score: 0.7885714285714286

```
In [68]: accuracy_score(ytest, y_best_rf_preds)
```

Out[68]: 0.8114285714285714


```
In [69]: y_best_rf_probas = best_rf_clf.predict_proba(Xtest)[: ,1]
```

```
In [70]: print(classification_report(ytest, (y_best_rf_probas > 0.5).astype(int)  
))
```

	precision	recall	f1-score	support
0	0.84	0.93	0.88	132
1	0.68	0.44	0.54	43
accuracy			0.81	175
macro avg	0.76	0.69	0.71	175
weighted avg	0.80	0.81	0.80	175

```
In [71]: fig, ax = plt.subplots(figsize=(8,4))  
features = train.columns  
importances = best_rf_clf.feature_importances_  
indices = np.argsort(importances)  
  
plt.title('Feature Importances')  
plt.barh(range(len(indices)), importances[indices], color='b', align='center')
```

```
plt.yticks(range(len(indices)), [features[i] for i in indices])
plt.xlabel('Relative Importance')
plt.show()
```



```
In [72]: y_rf_probs = clf_rf.predict_proba(Xtest)
```

```
In [73]: precision_rf, recall_rf, thresholds_rf = precision_recall_curve(ytest,
y_best_rf_probas)
```

```
In [74]: fpr_rf, tpr_rf, thresholds_rf = roc_curve(ytest, y_best_rf_probas)
```

```
In [ ]: #####logistic regression#####
basic approach to classification in supervised learning. Assumptions =
that data has no outliers, there are two classes to be predicted, and t
hat no two independent variables are highly correlated to each other.
```

```
In [75]: clf_log = LogisticRegression()
```

```
In [76]: train_scale = scale(train)
```

```
In [77]: Xtrain, Xtest, ytrain, ytest = train_test_split(X, y, random_state=42,
test_size=0.25)
```

```
In [78]: clf_log.fit(Xtrain, ytrain)
```

```
C:\Users\chinnababu\Anaconda3\lib\site-packages\sklearn\linear_model\lo
gistic.py:432: FutureWarning: Default solver will be changed to 'lbfgs'
in 0.22. Specify a solver to silence this warning.
FutureWarning)
```

```
Out[78]: LogisticRegression(C=1.0, class_weight=None, dual=False, fit_intercept=
True,
                                intercept_scaling=1, l1_ratio=None, max_iter=100,
                                multi_class='warn', n_jobs=None, penalty='l2',
                                random_state=None, solver='warn', tol=0.0001, verbos
e=0,
                                warm_start=False)
```

```
In [79]: y_log_pred = clf_log.predict(Xtest)
```

```
In [80]: accuracy_score(ytest, y_log_pred)
```

```
Out[80]: 0.8571428571428571
```

```
In [81]: C_space = np.array([0.0001, 0.001, 0.1, 1])
```

```
In [82]: param_grid = {'C': C_space}
```

```
In [83]: clf_log_tuning = GridSearchCV(clf_log, param_grid, cv=5)
```

```
In [84]: clf_log_tuning.fit(Xtrain, ytrain)
```

```
C:\Users\chinnababu\Anaconda3\lib\site-packages\sklearn\linear_model\lo
gistic.py:432: FutureWarning: Default solver will be changed to 'lbfgs'
in 0.22. Specify a solver to silence this warning.
FutureWarning)
```

[illegible]

```
FutureWarning)
C:\Users\chinnababu\Anaconda3\lib\site-packages\sklearn\linear_model\lo
gistic.py:432: FutureWarning: Default solver will be changed to 'lbfgs'
in 0.22. Specify a solver to silence this warning.
FutureWarning)
C:\Users\chinnababu\Anaconda3\lib\site-packages\sklearn\linear_model\lo
gistic.py:432: FutureWarning: Default solver will be changed to 'lbfgs'
in 0.22. Specify a solver to silence this warning.
FutureWarning)
C:\Users\chinnababu\Anaconda3\lib\site-packages\sklearn\linear_model\lo
gistic.py:432: FutureWarning: Default solver will be changed to 'lbfgs'
in 0.22. Specify a solver to silence this warning.
FutureWarning)
C:\Users\chinnababu\Anaconda3\lib\site-packages\sklearn\linear_model\lo
gistic.py:432: FutureWarning: Default solver will be changed to 'lbfgs'
in 0.22. Specify a solver to silence this warning.
FutureWarning)
C:\Users\chinnababu\Anaconda3\lib\site-packages\sklearn\linear_model\lo
gistic.py:432: FutureWarning: Default solver will be changed to 'lbfgs'
in 0.22. Specify a solver to silence this warning.
FutureWarning)
C:\Users\chinnababu\Anaconda3\lib\site-packages\sklearn\linear_model\lo
gistic.py:432: FutureWarning: Default solver will be changed to 'lbfgs'
in 0.22. Specify a solver to silence this warning.
FutureWarning)
C:\Users\chinnababu\Anaconda3\lib\site-packages\sklearn\linear_model\lo
gistic.py:432: FutureWarning: Default solver will be changed to 'lbfgs'
in 0.22. Specify a solver to silence this warning.
FutureWarning)
C:\Users\chinnababu\Anaconda3\lib\site-packages\sklearn\linear_model\lo
gistic.py:432: FutureWarning: Default solver will be changed to 'lbfgs'
in 0.22. Specify a solver to silence this warning.
FutureWarning)
C:\Users\chinnababu\Anaconda3\lib\site-packages\sklearn\linear_model\lo
gistic.py:432: FutureWarning: Default solver will be changed to 'lbfgs'
in 0.22. Specify a solver to silence this warning.
FutureWarning)
C:\Users\chinnababu\Anaconda3\lib\site-packages\sklearn\linear_model\lo
gistic.py:432: FutureWarning: Default solver will be changed to 'lbfgs'
```

```
in 0.22. Specify a solver to silence this warning.  
FutureWarning)
```

```
Out[84]: GridSearchCV(cv=5, error_score='raise-deprecating',  
                      estimator=LogisticRegression(C=1.0, class_weight=None, dual=False,  
                                                    fit_intercept=True,  
                                                    intercept_scaling=1, l1_ratio=None,  
                                                    max_iter=100, multi_class='warn',  
                                                    n_jobs=None, penalty='l2',  
                                                    random_state=None, solver='warn',  
                                                    tol=0.0001, verbose=0, warm_start=False),  
                      iid='warn', n_jobs=None,  
                      param_grid={'C': array([1.e-04, 1.e-03, 1.e-01, 1.e+00])},  
                      pre_dispatch='2*n_jobs', refit=True, return_train_score=False,  
                      scoring=None, verbose=0)
```

```
In [85]: print("Tuned Logistic Regression Parameters: {}".format(clf_log_tuning  
                      .best_params_))  
         print("Best score is {}".format(clf_log_tuning .best_score_))
```

```
Tuned Logistic Regression Parameters: {'C': 1.0}  
Best score is 0.7904761904761904
```

```
In [86]: clf_log = LogisticRegression(C = 1.0)
```

```
In [87]: clf_log.fit(Xtrain, ytrain)
```

```
C:\Users\chinnababu\Anaconda3\lib\site-packages\sklearn\linear_model\logistic.py:432: FutureWarning: Default solver will be changed to 'lbfgs'  
in 0.22. Specify a solver to silence this warning.  
FutureWarning)
```

```
Out[87]: LogisticRegression(C=1.0, class_weight=None, dual=False, fit_intercept=
```



```

True,
                                intercept_scaling=1, ll_ratio=None, max_iter=100,
                                multi_class='warn', n_jobs=None, penalty='l2',
                                random_state=None, solver='warn', tol=0.0001, verbos
e=0,
                                warm_start=False)

```

```
In [88]: y_preds = clf_log.predict(Xtest)
```

```
In [89]: p_clf_log_ba = clf_log.predict_proba(Xtest)
```

```
In [90]: accuracy_score(ytest, y_preds)
```

```
Out[90]: 0.8571428571428571
```

```
In [91]: print(classification_report(ytest, y_preds))
```

	precision	recall	f1-score	support
0	0.87	0.95	0.91	132
1	0.78	0.58	0.67	43
accuracy			0.86	175
macro avg	0.83	0.76	0.79	175
weighted avg	0.85	0.86	0.85	175


```
In [92]: precision_lg, recall_lg, thresholds_lg = precision_recall_curve(ytest,
p_clf_log_ba[:, 1])
```

```
In [93]: fpr_lg, tpr_lg, thresholds_lg = roc_curve(ytest, p_clf_log_ba[:, 1])
```

```
In [94]: fig, ax = plt.subplots(figsize=(8,8))
plt.plot(recall_lg, precision_lg)
plt.plot(recall_rf, precision_rf)
plt.plot(recall_dc, precision_dc)
plt.legend(('Logistic Reg', 'RandomForestClassifier', 'Decision Tree'))
```

```
plt.xlabel('Recall')
plt.ylabel('Precision')
plt.title('Precision vs. Recall curve')
```

Out[94]: Text(0.5, 1.0, 'Precision vs. Recall curve')




```
In [95]: area_log_reg = auc(recall_lg, precision_lg)
print(area_log_reg)
area_rf = auc(recall_rf, precision_rf)
```

```
print(area_rf)
area_dc = auc(recall_dc, precision_dc)
print(area_dc)
```

```
0.8011770322280233
0.6231740234944497
0.6083448382291768
```

```
In [96]: fig, ax = plt.subplots(figsize=(8,8))
plt.plot(fpr_rf, tpr_rf)
plt.plot(fpr_lg, tpr_lg)
plt.plot(fpr_dc, tpr_dc)
plt.legend(('RandomForestClassifier', 'Logistic Regression' , 'Decision
tree'))
plt.xlabel('False Positive Rate')
plt.ylabel('True Positive Rate')
plt.title('ROC Curves')
```

```
Out[96]: Text(0.5, 1.0, 'ROC Curves')
```



```
In [97]: Areas_ROC_decision = roc_auc_score(ytest, dt_pred1)
Areas_ROC_logistic = roc_auc_score(ytest, p_clf_log_ba[:, 1])
Areas_ROC_randomforest = roc_auc_score(ytest, y_best_rf_probas)
print(Areas_ROC_decision)
print(Areas_ROC_logistic)
print(Areas_ROC_randomforest)
```

0.7791578576462297

```
0.8879492600422833  
0.8095489781536294
```

In []: